



**BITS Pilani**

Pilani | Dubai | Goa | Hyderabad

# Data Structures and Algorithms

## CS F211

**Vishal Gupta**

Department of Computer Science and Information Systems  
Birla Institute of Technology and Science  
Pilani Campus, Pilani



## Agenda: Red-Black Trees

# Red-Black Trees

---

- *Red-black trees*:
  - Binary search trees augmented with node color
  - Operations designed to guarantee that the height  $h = O(\lg n)$



# Red-Black Properties

---

- The *red-black properties*:
  1. Every node is either red or black
  2. Every leaf (NULL pointer) is black
    - Note: this means every “real” node has 2 children
  3. If a node is red, both children are black
    - Note: can’t have 2 consecutive reds on a path
  4. Every path from node to descendent leaf contains the same number of black nodes
  5. The root is always black

# Red-Black Tree


- Each node has the following attributes:
  - Color
  - Key
  - Left child
  - Right child
  - Parent
- If a child of the parent of a node does not exist, the corresponding pointer attribute of the node contains the value NIL.
- We shall regard these NIL's as being pointers to leaves (external nodes) of the BST and the normal, key bearing nodes as being internal nodes of the tree.
- By constraining the node colors on any simple path from the root to a leaf, red black trees ensure that no such path is more than twice as long as any other path.



# Black-Height

---

- *Black-height*: The black-height of a node,  $X$ , in a red-black tree is the number of Black nodes on any path to a NULL (or leaf), not counting  $X$ .




## A Red-Black Tree with NULLs shown

Black-Height of the tree (the root) = 3

Black-Height of node “X” = 2






A Red-Black Tree with  
Black-Height = \_\_\_\_\_

# Red-Black Trees: An Example

- *Color this tree:*




Red-black properties:

1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# Red-Black Trees:

## The Problem With Insertion

- Insert 8
  - *Where does it go?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# Red-Black Trees:

## The Problem With Insertion

- Insert 8
  - *Where does it go?*
  - *What color should it be?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# Red-Black Trees:

## The Problem With Insertion

- Insert 8
  - *Where does it go?*
  - *What color should it be?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# Red-Black Trees:

## The Problem With Insertion

- Insert 11
  - *Where does it go?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# Red-Black Trees:

## The Problem With Insertion

- Insert 11
  - *Where does it go?*
  - *What color?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# Red-Black Trees:

## The Problem With Insertion

- Insert 11
  - *Where does it go?*
  - *What color?*
    - Can't be red! (#3)




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black



# Red-Black Trees:

## The Problem With Insertion

- Insert 11
  - *Where does it go?*
  - *What color?*
    - Can't be red! (#3)
    - Can't be black! (#4)




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# Red-Black Trees:

## The Problem With Insertion

- Insert 11
  - *Where does it go?*
  - *What color?*
    - Solution:  
recolor the tree




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# Red-Black Trees:

## The Problem With Insertion

- Insert 10
  - *Where does it go?*

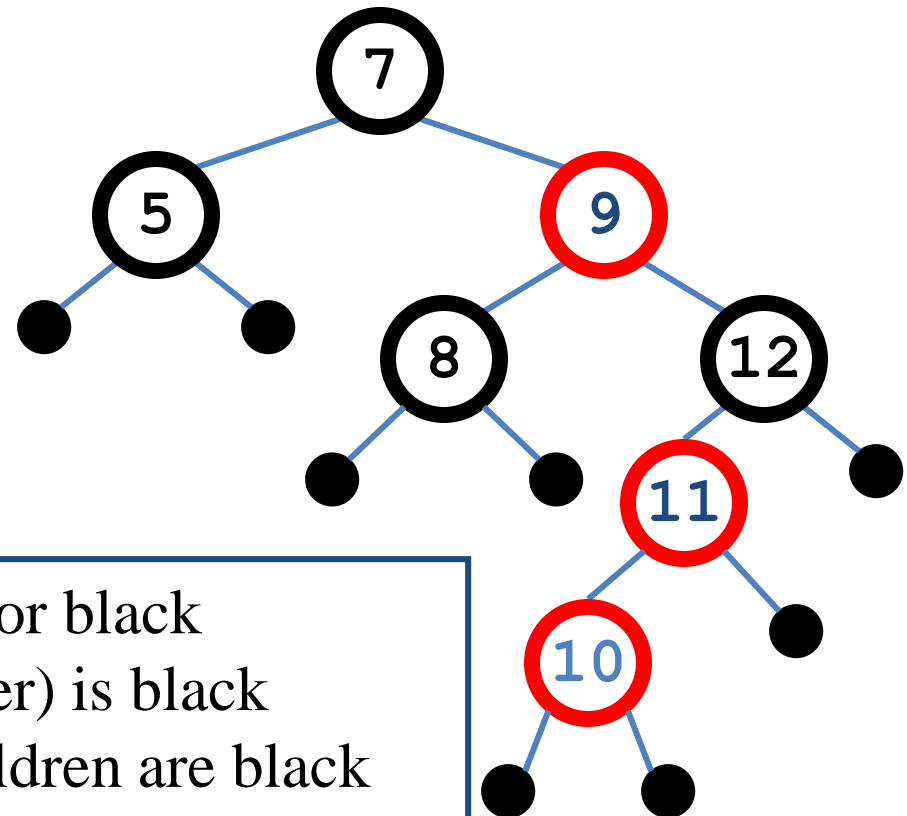


1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# Red-Black Trees:

## The Problem With Insertion

- Insert 10
  - *Where does it go?*
  - *What color?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

# Red-Black Trees:


## The Problem With Insertion

- Insert 10
  - *Where does it go?*
  - *What color?*
    - A: no color! Tree is too imbalanced
    - Must change tree structure to allow recoloring
  - Goal: restructure tree in  $O(\lg n)$  time




# RB Trees: Rotation

- Our basic operation for changing tree structure is called *rotation*:



- Does rotation preserve inorder key ordering?*




**Left-Rotate(T,x)**

```

y = right(x)           ; assume right(x) <> NIL
right(x) = left(y)      ; move y's child over
if left(y) <> NIL
then parent(left(y)) = x
parent(y) = parent(x)   ; move y up to x's position
if parent(x) = NIL
then root(T) = y
else if x = left(parent(x))
    then left(parent(x)) = y
    else right(parent(x)) = y
left(y) = x             ; move x down
parent(x) = y

```



**Right-Rotate(T,y)**

**`x = left(y)` ; assume `left(y) <> NIL`**

**`left(y) = right(x)`**

**if `right(x) <> NIL`**

**then `parent(right(x)) = y`**

**`parent(x) = parent(y)`**

**if `parent(y) = NIL`**

**then `root(T) = x`**

**else if `y = left(parent(y))`**

**then `left(parent(y)) = x`**

**else `right(parent(y)) = x`**


**`right(x) = y`**

**`parent(y) = x`**




# Rotation Example

- Rotate left about 9:



# Rotation Example

- Rotate left about 9:



# Red-Black Trees: Insertion

---

- Insertion: the basic idea
  - Insert  $x$  into tree, color  $x$  red
  - Which of the red-black properties might be violated?
    - Root is always black
    - Red node cannot have a red child
  - Fix the violated properties.

# Red-Black Trees: Insertion


---

## Insertion

1. Insert node into tree using BST Insert( $T, x$ ) and color node Red
2. Fix violated RBT properties
  1. Root is always black
  2. Red node cannot have a red child
3. Color root Black

# Red-Black Trees: Insertion

- If parent node 'a' was Black, then no changes are necessary.
- If not, then there are following cases to consider for each of the orientations below.



- Move up the tree until there are no violations or we are at the root.
- In the following discussion we will assume the parent is a left child (if the parent is a right child perform the same steps swapping ``right" and ``left")

# Red-Black Trees: Insertion


RB-Insert( $T, x$ )

Case I:  $x$ 's uncle is Red

- Change  $x$ 's grandparent to Red
- Change  $x$ 's uncle and parent to Black
- Change  $x$  to  $x$ 's grandparent

**How to get uncle ( $x$ )**

if  $\text{parent}(x) = \text{left}(\text{parent}(\text{parent}(x)))$   
 then  $\text{uncle}(x) = \text{right}(\text{parent}(\text{parent}(x)))$   
 else  $\text{uncle}(x) = \text{left}(\text{parent}(\text{parent}(x)))$



# Red-Black Trees: Insertion

RB-Insert( $T, x$ )

Case II:  $x$ 's uncle is Black,  $x$  is the right child of its parent

- Change  $x$  to  $x$ 's parent
- Rotate  $x$ 's parent (now  $x$ ) left to make Case III
- Case II is now Case III

Case III:  $x$ 's uncle is Black,  $x$  is the left child of its parent

- Set  $x$ 's parent to Black
- Set  $x$ 's grandparent to Red
- Rotate  $x$ 's grandparent right

# Red-Black Trees: Insertion


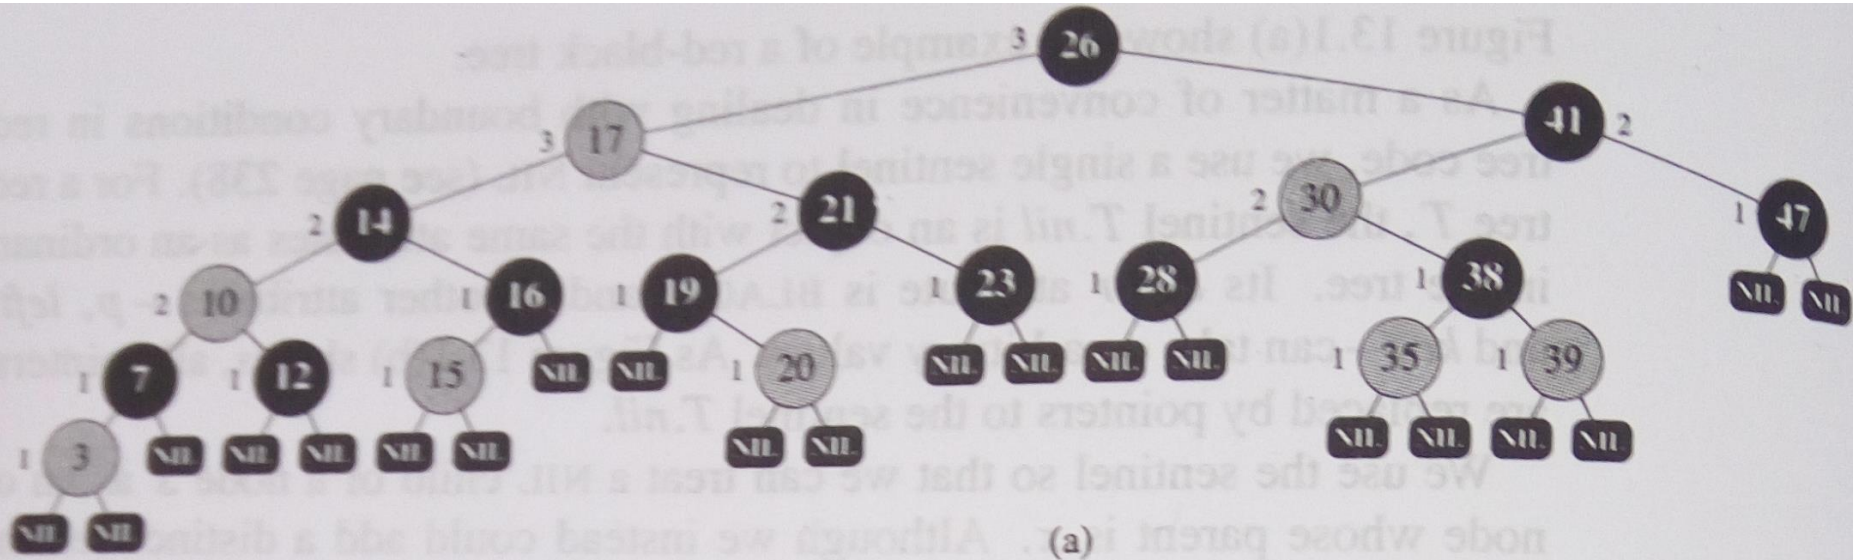
---

Example:

Insert the following keys in a Red Black Tree in order:

3   2   5   6   9   4   7   8





---

**Theorem 1** – In a red-black tree, any subtree rooted at  $x$  contains at least  $(2^{bh(x)} - 1)$  internal nodes, where  $bh(x)$  is the black height of node  $x$ .

Proof: by induction on height of  $x$ .

In a red-black tree, at least half the nodes on any path from the root to a NULL (i.e. leaf) must be Black.

**Proof** – If there is a Red node on the path, there must be a corresponding Black node.

Algebraically this theorem means

$$bh(x) \geq h/2$$

where  $x$  is the root node.

In a red-black tree, no path from any node, X, to a NULL (i.e. leaf) is more than twice as long as any other path from X to any other NULL (i.e. leaf).

Proof: By definition, every path from a node to any NULL contains the same number of Black nodes. By Theorem 2, at least  $\frac{1}{2}$  the nodes on any such path are Black. Therefore, there can be no more than twice as many nodes on any path from X to a NULL as on any other path. Therefore the length of every path is no more than twice as long as any other path.

**Theorem 4** – A red-black tree with  $n$  internal nodes has height  $h \leq 2 \lg(n + 1)$ .

**Proof:** Let  $h$  be the height of the red-black tree with root  $x$ .  
By Theorem 2,

$$bh(x) \geq h/2$$

From Theorem 1,  $n \geq 2^{bh(x)} - 1$

Therefore  $n \geq 2^{h/2} - 1$

$$n + 1 \geq 2^{h/2}$$

$$\lg(n + 1) \geq h/2$$

$$2\lg(n + 1) \geq h$$

# RB Trees: Worst-Case Time

---

- So we've proved that a red-black tree has  $O(\lg n)$  height
- Corollary: These operations take  $O(\lg n)$  time:
  - Minimum(), Maximum()
  - Successor(), Predecessor()
  - Search()
- Insert() and Delete():
  - Will also take  $O(\lg n)$  time
  - But will need special care since they modify tree



# Red-Black Trees: Deletion

---

- As insert had three cases, delete has four different cases.
- Do it yourself.



**BITS Pilani**

Pilani | Dubai | Goa | Hyderabad



**Thank You**