



BITS Pilani

Pilani | Dubai | Goa | Hyderabad

Data Structures and Algorithms

CS F211

Vishal Gupta

Department of Computer Science and Information Systems
Birla Institute of Technology and Science
Pilani Campus, Pilani



Agenda: B Trees



B Tree: Motivation

Motivation for B-Trees

- Index structures for large datasets cannot be stored in main memory
- Storing it on disk requires different approach to efficiency
- Assuming that a disk spins at 3600 RPM, one revolution occurs in $1/60$ of a second, or 16.7ms
- Crudely speaking, one disk access takes about the same time as 200,000 instructions

Motivation (cont.)

- Assume that we use an AVL tree to store about 20 million records
- We end up with a **very** deep binary tree with lots of different disk accesses; $\log_2 20,000,000$ is about 24, so this takes about 0.2 seconds
- We know we can't improve on the $\log n$ lower bound on search for a binary tree
- But, the solution is to use more branches and thus reduce the height of the tree!
 - As branching increases, depth decreases

B-Trees




B-Trees are useful in the following cases:


- The number of objects is too large to fit in memory.
- Need external storage.
- Disk accesses are slow, thus need to minimize the number of disk accesses

Is RB Tree good in these situations?

B Tree Example



A B-tree of height 2 containing over one billion keys.



1 node,
1000 keys


1001 nodes,
1,001,000 keys

1,002,001 nodes,
1,002,001,000 keys

B-Trees



- B-Trees are balanced, like RB trees.
- They have a large number of children (large branching factor), unlike RB trees.
- The branching factor is determined by the size of disk transfers (page size).
- Each object (node) referenced requires a DiskRead.
- Each object modified requires a DiskWrite.
- The root of the tree is kept in memory at all times.
- Insert, Delete, Search = $O(h)$, where h is the height of the tree.
 $O(\lg n)$, though much less in reality ($\log_{BF} n$).




Properties of B-Trees

- Every leaf has the same depth equal to the height of the tree.
- The number of keys is bounded in terms of the minimum degree $t \geq 2$. (WHY?)
- $n(x) \geq t-1$ (except root ≥ 1)
- $\#children(x) \geq t$ (except root ≥ 0), leaves = 0
- $n(x) \leq 2t - 1$
- $\#children \leq 2t$ (except leaves which = 0)
- If $n(x) = 2t - 1$ then n is a full node.

What is the height of B-Tree

- Given a B-Tree of height h , minimum degree t , and number of keys n , prove that the height of the B-tree is:

$$h \leq \log_t \frac{n+1}{2}$$



What is the height of B-Tree

- Given a B-Tree of height h , minimum degree t , and number of keys n , prove that the height of the B-tree is:

$$h \leq \log_t \frac{n+1}{2}$$


Searching a key



```
Search(x, k)
  if k in node x
  then return x and i such that  $\text{key}_i(x) = k$ 
  else if x is a leaf
    then return NIL
  else find i such that  $\text{key}_{i-1}(x) < k < \text{key}_i(x)$ 
    DiskRead(childi(x))
    return Search(childi(x), k)
```

What is the Time-Complexity ??

Search For: N



Inserting a Node



Overview

- If node x is a non-full ($< 2t-1$ keys) leaf, then insert new key k in node x
- If node x is non-full but not a leaf, then recurse to appropriate child of x
- If node x is full ($2t-1$ keys), then ``split'' the node into x_1 and x_2 , and recurse to appropriate node x_1 or x_2 .

Splitting a Node



B-Tree-Split-Child(x, i, y) ; x is parent, y is child in i th subtree
Allocate(z) ; $n(z)=t-1, \text{leaf}(z) = \text{leaf}(y)$
Copy y 's second half keys and children to z
 $n(y) = t-1$
Shift x 's keys and children one to the right from i
 $\text{child}_{i+1}(x) = z$
 $\text{key}_i(x) = \text{key}_t(y)$
 $n(x) = n(x) + 1$
Write(x)
DiskWrite(y)
DiskWrite(z)

Inserting a Node



Insert: B-Tree-Insert(T, k)

- Start at root(T) moving down the tree looking for the proper leaf to put k
- Split all full nodes along the way

Inserting a Node



B-Tree-Insert(T, k)

$r = \text{root}(T)$

 if $n(r) = 2t-1$; full

 then allocate empty node s pointing to r

 B-Tree-Split-Child($s, 1, r$)

 B-Tree-Insert-Nonfull(s, k)

 else B-Tree-Insert-Nonfull(r, k)

B-Tree-Insert-Nonfull(x, k)

 if leaf(x)

 then shift keys of x higher than k one to the right

 put k in appropriate spot

$n(x) = n(x) + 1$

 DiskWrite(x)


 else find smallest i such that $k < \text{key}_i(x)$

 DiskRead(child $_i(x)$)

 if $n(\text{child}_i(x)) = 2t - 1$; full


 then B-Tree-Split-Child($x, i, \text{child}_i(x)$)

Insertion Example




Insert B:

Insertion Example




Insert Q:

Insertion Example




Insert L:

Insertion Example



Insert F:

Insertion Example



Insert F:

Deleting a Node



Overview

Deletion: B-Tree-Delete(x , k)

- 1) Search down the tree for node containing k
- 2) When B-Tree-Delete is called recursively, the number of keys in x must be at least the minimum degree t (the root can have $< t$ keys)
 - 1) If x is a leaf, just remove key k and still have at least $t-1$ keys in x
 - 2) If there are not $\geq t$ keys in x , then borrow keys from other nodes.

Deleting a Node




Deletion

There are three general cases:

[Case 1:] If key k in node x and x is a leaf, then remove k from x .

Delete(x , 8)



[Case 2:] If k is in x and x is an internal node.

One of three subcases:

Deleting a Node




Case 2a

If child y preceding k in x has $\geq t$ keys:

- Find predecessor k' of k in subtree rooted at y
- Recursively delete k' (first two steps can be performed in one pass down the tree)
- Replace k by k' in x

Delete(x , 8)



Deleting a Node




Case 2b

If child z following k in x has $\geq t$ keys:

- Find successor k' of k in subtree y
- Recursively delete k'
- Replace k by k' in x

Delete(x, 4)




Deleting a Node



Case 2c

If both y and z have $t-1$ keys:

- Merge k and all of z into y
- Free z
- Recursively delete k from y



Deleting a Node



Case 3

if k not in internal node x

then determine subtree $\text{child}_i(x)$ containing k

if $\text{child}_i(x)$ has $\geq t$ keys

then $\text{B-Tree-Delete}(\text{child}_i(x), k)$

else execute Case 3a or 3b until can descend to node having $\geq t$ keys

Deleting a Node




Case 3a

If $\text{child}_i(x)$ has $t-1$ keys but has a left or right sibling with $\geq t$ keys, then borrow one from sibling

- move key from x to $\text{child}_i(x)$
- move key from sibling to x
- move child from sibling to $\text{child}_i(x)$

Delete($x, 1$)




Deleting a Node



Case 3b

If $\text{child}_i(x)$ and its left and right siblings have $t-1$ keys then merge $\text{child}_i(x)$ with one sibling using median key from x .

Delete(x, 11)



Deleting a Node




Deletion Example

1. If the key k is in node x and x is a leaf, delete the key k from x .
2. If the key k is in node x and x is an internal node, do the following.
 - a) If the child y that precedes k in node x has at least t keys, then find the predecessor k' of k in the subtree rooted at y . Recursively delete k' , and replace k by k' in x . (Finding k' and deleting it can be performed in a single downward pass.)
 - b) Symmetrically, if the child z that follows k in node x has at least t keys, then find the successor k' of k in the subtree rooted at z . Recursively delete k' , and replace k by k' in x . (Finding k' and deleting it can be performed in a single downward pass.)
 - a) Otherwise, if both y and z have only $t - 1$ keys, merge k and all of z into y , so that x loses both k and the pointer to z , and y now contains $2t - 1$ keys. Then, free z and recursively delete k from y .

3. If the key k is not present in internal node x , determine the root $c_i[x]$ of the appropriate subtree that must contain k , if k is in the tree at all. If $c_i[x]$ has only $t - 1$ keys, execute step 3a or 3b as necessary to guarantee that we descend to a node containing at least t keys. Then, finish by recursing on the appropriate child of x .


- a) If $c_i[x]$ has only $t - 1$ keys but has a sibling with t keys, give $c_i[x]$ an extra key by moving a key from x down into $c_i[x]$, moving a key from $c_i[x]$'s immediate left or right sibling up into x , and moving the appropriate child from the sibling into $c_i[x]$.
- a) If $c_i[x]$ and all of $c_i[x]$'s siblings have $t - 1$ keys, merge c_i with one sibling, which involves moving a key from x down into the new merged node to become the median key for that node.

Deletion Example



Delete F


Deletion Example



Delete M


Deletion Example

(c) *M* deleted: case 2a




Delete G

Deletion Example




Delete D

Deletion Example



Delete B

Deletion Example



Reasons for using B-Trees

- When searching tables held on disc, the cost of each disc transfer is high but doesn't depend much on the amount of data transferred, especially if consecutive items are transferred
 - If we use a B-tree of order 101, say, we can transfer each node in one disc read operation
 - A B-tree of order 101 and height 3 can hold $101^4 - 1$ items (approximately 100 million) and any item can be accessed with 3 disc reads (assuming we hold the root in memory)
- If we take $m = 3$, we get a **2-3 tree**, in which non-leaf nodes have two or three children (i.e., one or two keys)
 - B-Trees are always balanced (since the leaves are all at the same level), so 2-3 trees make a good type of balanced tree

Comparing Trees

- Binary trees
 - Can become *unbalanced* and *lose* their good time complexity (big O)
 - AVL trees are strict binary trees that *overcome the balance problem*
 - Heaps remain balanced but only *prioritise* (not order) the keys
- Multi-way trees
 - B-Trees can be *m*-way, they can have any (odd) number of children
 - One B-Tree, the 2-3 (or 3-way) B-Tree, *approximates* a permanently balanced binary tree, exchanging the AVL tree's balancing operations for insertion and (more complex) deletion operations



BITS Pilani

Pilani | Dubai | Goa | Hyderabad



Thank You